

Introduction to Data Science and Machine Learning

Research Assignment

Question 1: Defining Data Science and Machine Learning

1.1 What is Data Science?

Data Science is an evidence-driven discipline that combines statistical reasoning, computing, and domain expertise to convert raw information into practical knowledge for decision-making. Rather than focusing only on analysis, it manages the entire process from identifying a problem through preparing data, interpreting results, and delivering actionable solutions.

1.2 What is Machine Learning?

Machine Learning is a computational approach in which algorithms infer patterns from historical data and improve their performance through experience instead of relying solely on manually coded rules. Its primary value lies in prediction, classification, and automated pattern discovery. ML algorithms rely on statistical models and iterative optimization processes.

1.3 The Relationship Between Data Science and Machine Learning

Data Science and Machine Learning are complementary but not interchangeable. **Data Science** provides the broader analytical framework for framing questions, collecting and preparing data, interpreting outcomes, and communicating findings, **while Machine Learning** supplies predictive techniques that can be embedded within that workflow when automation or forecasting is required.

Pangeanic (2023) captures this well: within the entire analysis process performed by data science, machine learning provides the techniques and tools necessary for building algorithms or models that learn by extracting meaning from the data studied. Their common meeting point is data itself.

Aspect	Data Science	Machine Learning
Scope	Broad end-to-end discipline	Focused on learning algorithms
Goal	Extract insights and inform decisions	Build predictive or pattern-finding models
Includes	Statistics, visualization, engineering, ML	Mathematical models, optimization
Human involvement	High (storytelling, strategy)	Moderate (algorithm selection, tuning)
Output	Insights, dashboards, recommendations	Trained models and predictions

1.4 Real-Life Example: Credit Card Fraud Detection

A compelling real-world illustration of how Data Science and Machine Learning work together is credit card fraud detection. As described by Syracuse University (2025), while Data Science provides the statistical foundation by structuring and analyzing financial records, Machine Learning enhances the actual process of fraud detection by recognizing subtle patterns that might indicate suspicious behavior.

In practice, data scientists first collect and clean millions of transaction records (Data Science phase). They explore patterns, identify useful features such as transaction time, location, and amount, and then feed the labeled historical data (fraud vs. not fraud) into a Machine Learning classifier. The trained model can then detect fraudulent transactions in real time without any human checking each transaction manually. This demonstrates the synergy: Data Science sets up the problem and the data; Machine Learning automates the pattern recognition at scale.

Question 2: The Data Science Lifecycle and Where ML Fits

2.1 The Data Science Lifecycle

A data science project does not begin with modeling. It follows a structured lifecycle with distinct stages, each building on the previous one. Skipping any stage introduces risk of failure or misleading results. The lifecycle can be described as follows:

Stage	Description	Key Activities
1. Problem Definition	Identify and frame the question to be solved	Stakeholder interviews, goal setting, success metric definition
2. Data Collection	Gather raw data from relevant sources	Databases, APIs, sensors, public datasets, web scraping
3. Data Cleaning	Prepare and preprocess raw data for analysis	Handle missing values, remove duplicates, normalize formats
4. Exploratory Data Analysis (EDA)	Understand the structure and distribution of data	Statistical summaries, correlations, visualizations
5. Feature Engineering	Transform raw variables into informative model inputs	Encoding, scaling, creating derived variables
6. Modeling (ML)	Train a predictive or classification model	Algorithm selection, training, hyperparameter tuning
7. Evaluation	Assess model performance on unseen data	Accuracy, precision, recall, MSE, cross-validation
8. Deployment	Integrate the model into production systems	APIs, dashboards, mobile apps, monitoring pipelines

2.2 Where Does Machine Learning Fit?

Machine Learning fits primarily at Stage 6 (Modeling), but its influence extends to several surrounding stages. Feature Engineering (Stage 5) is directly motivated by what the chosen ML algorithm needs. Evaluation (Stage 7) is shaped by how the ML model performs. Deployment (Stage 8) is determined by the nature of the trained model.

Importantly, not every data science project reaches the modeling stage.

sometimes the insight from exploratory data analysis alone is sufficient.

ML is introduced when the objective requires prediction, classification, or pattern discovery at a scale and complexity that rules-based approaches cannot handle.

The reason ML belongs in the middle of the lifecycle rather than at the beginning is that it absolutely depends on what comes before it. A model trained on dirty, poorly understood, or unrepresentative data will fail regardless of how sophisticated the algorithm is.

This is why practitioners often say that data preparation consumes 70-80% of a project's effort.

Question 3: Supervised vs. Unsupervised Learning

3.1 Supervised Learning

Supervised learning trains a model with examples that already include the desired answer, enabling the algorithm to map inputs to known outputs and generalize that mapping to future observations.

Supervised learning is used for two main task types: **regression** (predicting a continuous numerical value) and **classification** (assigning inputs to discrete categories). A well-known example is email spam detection. A classifier is trained on thousands of emails previously labeled as 'spam' or 'not spam.' Once trained, it can automatically categorize new emails it has never seen before by recognizing learned patterns in the text, links, and sender metadata.

3.2 Unsupervised Learning

Unsupervised learning analyzes datasets that lack predefined target labels, seeking to reveal latent structure.

A practical example is customer segmentation. An e-commerce company may have purchase history data for millions of customers but no pre-assigned labels for customer type. An unsupervised clustering algorithm such as K-Means can group customers into segments based on purchasing behavior without any human labeling. These discovered segments can then inform targeted marketing campaigns.

3.3 Comparison Table

Dimension	Supervised Learning	Unsupervised Learning
Data requirement	Labeled (input + correct output)	Unlabeled (input only)
Human involvement	High (labeling required upfront)	Low during training; higher for interpreting results
Goal	Predict or classify known outcomes	Discover hidden patterns or groupings
Common algorithms	Linear Regression, SVM, Neural Networks, Decision Trees	K-Means, DBSCAN, PCA, Autoencoders
Task types	Regression, Classification	Clustering, Dimensionality Reduction
Example use case	Spam filter, house price prediction	Customer segmentation, anomaly detection
Accuracy	Generally higher with good labels	Harder to evaluate objectively

Databricks (2025) provides a useful framing: **supervised learning** focuses on prediction and classification against known outcomes, while **unsupervised learning** focuses on discovery and insight generation. Neither approach is universally superior; the choice depends on whether labeled data is available and what the project goal is.

Question 4: Overfitting - Causes and Prevention

4.1 What is Overfitting?

Overfitting occurs when a model captures accidental characteristics of its training examples instead of the underlying signal, producing impressive training performance but weak results on genuinely new data.

The conceptual opposite, **underfitting**, occurs when a model is too simple to capture even the true patterns in the training data. Both sit at opposite ends of the bias-variance tradeoff: overfitting corresponds to high variance (sensitivity to training data specifics), while underfitting corresponds to high bias (systematic errors due to overly simplistic assumptions).

4.2 Causes of Overfitting

Several factors contribute to overfitting, as documented by Lightly.ai (2025) and Meegle (n.d.):

- **Insufficient training data:** When the dataset is small, a complex model can memorize all examples rather than learning generalizable rules.
- **Excessive model complexity:** Using a model with too many parameters or layers relative to the amount of training data gives it enough capacity to fit noise as well as signal.
- **Noise and irrelevant features:** If the dataset contain errors, outliers, or features that are statistically coincidental rather than causally related to the target, the model may learn false correlations.
- **Training for too many iterations:** In iterative algorithms, training beyond the point of convergence on meaningful patterns causes the model to start fitting residual noise.

4.3 Prevention Techniques

The machine learning literature documents several well-established techniques for mitigating overfitting (Towards Data Science, 2025; Number Analytics, 2025):

Technique	How It Works	When to Use
Train-Test Split / Hold-out	Reserve a portion of data for evaluation only; never train on it	Always; foundational practice
Cross-Validation (K-Fold)	Rotate which subset is used for testing across K iterations for more robust estimate	When data is limited
Regularization (L1/L2)	Add a penalty term to the loss function that discourages overly large model weights	Linear/logistic regression, neural networks
Dropout (Neural Networks)	Randomly deactivate neurons during training to prevent co-dependence	Deep learning models
Early Stopping	Halt training when validation performance stops improving	Iterative learning algorithms
Data Augmentation	Artificially increase training set size by transforming existing examples	Image, text, and audio data
Feature Selection	Remove irrelevant or redundant input features	High-dimensional datasets
Ensemble Methods	Combine multiple models to average out individual overfitting (e.g., Random Forest)	When single models are unstable

The underlying principle across all these techniques is the same: ensure the model is evaluated on data it has never seen during training, and constrain its complexity to what the evidence in the training data can genuinely support.

Question 5: Train-Test Split

5.1 What is Train-Test Split?

Train-test split is a fundamental model validation technique in machine learning that involves dividing a dataset into two non-overlapping subsets: a training set and a testing set. Built In (n.d.) defines it as a model validation technique that divides data into training and testing sets to simulate how models perform on new data. The training set is used to fit the model, while the test set provides an honest evaluation of the model's performance on data it has never encountered. A typical split ratio is 80% training and 20% testing, though ratios of 70/30 are also common. For projects with enough data, a three-way split (training, validation, test) is preferred, where the validation set guides hyperparameter tuning and the test set provides the final, unbiased performance estimate.

5.2 Why Is It Necessary?

The necessity of the train-test split stems from a fundamental challenge: a model that is evaluated on the same data it was trained on will almost always appear to perform better than it actually does in the real world. This is because the model has, to some extent, memorized the training examples.

The University of Illinois Data Science Discovery resource explains the intuition clearly: the test/train split simulates new data that has never been seen before, since the test set is data not found in the training set. Without this separation, accuracy measurements are artificially inflated and give a false sense of confidence in the model.

5.3 The Role of Validation Set

Many practitioners extend the two-way split to a three-way split: training, validation, and test. Lightly.ai (2025) explains each role: the training set is the largest portion, used to train the model and optimize its internal parameters; the validation set is used during model development to fine-tune hyperparameters and detect overfitting; and the test set provides the final, unbiased evaluation. Common ratios include 70/15/15 or 80/10/10.

This distinction matters because if the test set is used repeatedly to make tuning decisions, it effectively becomes part of the training process and loses its value as an independent benchmark.

Question 6: Case Study - Machine Learning in Healthcare

6.1 Source

Khan, S.I., Aishi, A.A., Sazzad, S., Sayduzzaman, M., & Band, S.S. (2024). Machine learning and deep learning-based approach in smart healthcare: Recent advances, applications, challenges and opportunities. *AIMS Public Health*, 11(1), 58-109. <https://doi.org/10.3934/publichealth.2024004>

6.2 Overview

This peer-reviewed survey paper, published in January 2024 in *AIMS Public Health*, provides a comprehensive examination of how ML and Deep Learning (DL) are being applied across the spectrum of modern healthcare. The authors conducted a systematic review of the most prominent journal and conference databases to identify cutting-edge progress in ML and DL-based analysis for smart healthcare systems.

6.3 Key Findings

The study identifies several major domains in which ML is producing measurable clinical impact:

- **Disease Prediction and Early Diagnosis:** ML models trained on patient records and medical imaging data are achieving diagnostic accuracy comparable to specialist physicians in domains such as cancer detection, diabetic retinopathy classification, and cardiovascular risk stratification.
- **Drug Discovery:** ML algorithms accelerate the identification of promising drug candidates by learning molecular property patterns from vast chemical databases, reducing the time and cost of traditional laboratory screening.
- **Medical Image Analysis:** Convolutional Neural Networks (CNNs), a class of supervised deep learning models, are applied to X-rays, MRIs, and CT scans to detect abnormalities with high sensitivity. These are trained on labeled image datasets annotated by radiologists.
- **Personalized Treatment Recommendations:** By learning from large electronic health record (EHR) datasets, ML systems can suggest individualized treatment plans based on a patient's profile, history, and genetic markers.
- **Remote Patient Monitoring:** IoT-connected wearable devices generate continuous health streams that ML models analyze in real time to flag anomalies, reducing hospital readmission rates.

6.4 Data Science Lifecycle Stages Covered

Mapping this case study to the Data Science lifecycle reveals coverage across multiple stages:

DS Lifecycle Stage	How It Appears in This Case Study
Problem Definition	Each application area (disease prediction, drug discovery, etc.) represents a clearly framed clinical problem with measurable success criteria such as diagnostic accuracy or early detection rate.
Data Collection	Patient EHRs, medical imaging archives, wearable sensor streams, and genomic databases serve as primary data sources. The diversity and volume of data collected is a central focus of the paper.
Data Cleaning & Preprocessing	The authors explicitly highlight preprocessing of medical images (normalization, augmentation) and handling of missing clinical values as critical preconditions for model performance.
Feature Engineering	Extracting relevant clinical features from raw images, lab results, and sensor readings is described as a key step in preparing inputs for ML models.
Modeling (Machine Learning)	Supervised learning models (CNNs for imaging, regression models for risk scores) and unsupervised models (clustering for patient subgroups) are extensively discussed and evaluated.

DS Lifecycle Stage	How It Appears in This Case Study
Evaluation	The paper reports standard clinical ML metrics including sensitivity, specificity, AUC-ROC, and F1-score, benchmarked against both prior baselines and human expert performance.
Deployment & Monitoring	The paper discusses deployment challenges including data privacy regulations, model interpretability requirements (explain ability for clinicians), and the need for continuous retraining as patient populations evolve.

6.5 Significance

This study demonstrates that ML in healthcare is not a single-stage intervention but a full Data Science pipeline problem. The quality of clinical outcomes depends as much on data governance, preprocessing rigor, and deployment infrastructure as on the sophistication of the ML algorithm itself. The paper also raises important ethical considerations around algorithmic fairness, noting that models trained on non-representative datasets risk producing biased recommendations that could disproportionately affect underserved patient populations.

References

- AWS (Amazon Web Services). (2024). Supervised vs. unsupervised learning. <https://aws.amazon.com/compare/the-difference-between-machine-learning-supervised-and-unsupervised/>
- Built In. (n.d.). Train test split: What it means and how to use it. <https://builtin.com/data-science/train-test-split>
- Databricks. (2025). Supervised vs. unsupervised learning: Understanding the differences and capabilities of each ML approach. <https://www.databricks.com/blog/supervised-vs-unsupervised-learning>
- IBM. (2025). Data science vs. machine learning. <https://www.ibm.com/think/topics/data-science-vs-machine-learning>
- IBM. (2026). Supervised vs. unsupervised learning. <https://www.ibm.com/think/topics/supervised-vs-unsupervised-learning>
- Khan, S.I., Aishi, A.A., Sazzad, S., Sayduzzaman, M., & Band, S.S. (2024). Machine learning and deep learning-based approach in smart healthcare: Recent advances, applications, challenges and opportunities. *AIMS Public Health*, 11(1), 58-109. <https://doi.org/10.3934/publichealth.2024004>
- Lightly.ai. (2025). Overfitting in machine learning: Causes, detection, and prevention. <https://www.lightly.ai/blog/overfitting>
- Lightly.ai. (2025). Train-test-validation split. <https://www.lightly.ai/blog/train-test-validation-split>
- Meegle. (n.d.). Overfitting prevention techniques. https://www.meegle.com/en_us/topics/overfitting/overfitting-prevention-techniques
- Number Analytics. (2025). Effective overfitting prevention techniques to boost model quality. <https://www.numberanalytics.com/blog/effective-overfitting-prevention-techniques-ml-model-quality>
- Northdoor. (2024). Data science vs. machine learning: Understanding the difference. <https://www.northdoor.co.uk/insight/blog/data-science-vs-machine-learning-understanding-the-difference/>
- Pangeanic. (2023). The relationship between data science and machine learning. <https://blog.pangeanic.com/trelationship-between-data-science-and-machine-learning>
- Syracuse University iSchool. (2025). Data science vs. machine learning: Key differences explained. <https://ischool.syracuse.edu/data-science-vs-machine-learning/>
- Towards Data Science. (2025). 8 simple techniques to prevent overfitting. <https://towardsdatascience.com/8-simple-techniques-to-prevent-overfitting-4d443da2ef7d/>
- University of Illinois Data Science Discovery. (n.d.). Test/train split. <https://discovery.cs.illinois.edu/learn/Towards-Machine-Learning/TestTrain-Split/>
- VinUniversity. (2026). Data science vs. machine learning: Scope and skillsets. <https://vinuni.edu.vn/data-science-vs-machine-learning/>